THE SI SENTINEL

Weekday briefing on builders shaping super intelligence

MONDAY, NOVEMBER 9, 2026

THE MORNING PICTURE

Starship goes for a double catch, Cybertruck makes way for Cybercab, and Richard Sutton reopens the scaling argument

SpaceX set Starship Flight 15 for as early as Friday, Nov. 13, and Elon Musk said it will try to catch both the Super Heavy booster and the ship, which would be the first time every part of an orbital rocket comes back intact. Over the weekend he described Starlink V4 laser links at 10 terabits per second, confirmed that the Cybertruck line at Gigafactory Texas winds down around the middle of next year to make room for Cybercab, and promised overnight fixes after Grok Bot subscribers complained about usage limits.

Richard Sutton, the reinforcement-learning pioneer, wrote that "brute force computation is not enough for intelligence," and François Chollet and Albert Gu agreed from different angles. Patrick Collison asked what the positive AI scenario actually looks like, Vitalik Buterin defended slowing the data-center buildout, and Arthur Mensch warned that a French copyright bill could push Mistral's compute abroad. Meta put Muse Wonder-trained models in front of every Muse user, and Nvidia's board opened an investigation into reports that Jensen Huang was targeted by a foreign intelligence effort.

I · STARSHIP AND ORBIT

Flight 15 will try to catch both stages

SpaceX's mission page targets Friday, Nov. 13, for the first attempt to return and catch both Super Heavy and Starship at Starbase. Musk spent the weekend on what comes next: Starlink V4 laser links, orbital compute and a bigger Raptor engine.

Nov 13
Earliest Flight 15 launch date; the 75-minute window opens at 5:45pm EST, pending regulatory approval

26
Starlink V3 satellites the ship is to deploy before returning to Starbase

~22.5 hr
Planned time from liftoff to ship catch in SpaceX's flight timeline; the booster catch comes at about 6.5 minutes

Elon Musk @elonmusk · CEO, SpaceX and Tesla

SpaceX released a short film, "Go For Orbit," on Friday morning about the first orbital Starship launch carrying advanced Starlink satellites. Musk quoted it at 10:47am EST:

> *"SpaceX will attempt to catch both the Super Heavy booster and Starship on the upcoming flight. If both are successful, it will mark the first time that all parts of an orbital rocket return to Earth intact. This is the critical breakthrough needed to make life multiplanetary."*
> ELON MUSK · @ELONMUSK · 10:47AM EST · POST

SpaceX's Flight 15 page says it will be the first flight to return both Starship and Super Heavy to the launch site for catch, and the first return attempt for a Super Heavy V3 booster; three V1 boosters were caught on earlier flights. The ship is to deploy 26 Starlink V3 satellites and return only if the flight team confirms that the systems needed for reentry, the landing burn and the tower approach are healthy. Next Spaceflight, quoting SpaceX, puts the ship at about 275 km for roughly 14 orbits over more than 22 hours. Musk quoted SpaceX's Saturday update with "Catching a giant rocket falling from space" at 11:54am EST. Replying to a post about the Friday-the-13th launch date, he wrote: "Fate does love irony, so perhaps the unluckiest day will actually be the luckiest."



The launch tower's arms catch a Super Heavy booster at Starbase, Texas, during Starship's fifth flight test on Oct. 13, 2024.

PHOTO: STEVE JURVETSON / WIKIMEDIA COMMONS · CC BY 2.0 (CROPPED)

Starlink V4 laser links @elonmusk

In an onstage interview with investor Ron Baron on Friday, Musk said, in a clip shared by Sawyer Merritt: "Starlink version 4 will have 10x the capability of V3 (which itself is 10x better than V2), especially with the long-distance laser links. Our next-gen laser links will be 10-terabits per second. 100x better. And we'll put 10 of them on." He went further at 2:32pm EST:

> *"~100 lasers per satellite, each operating above 10 Tb/s, with over 100k satellites should be able to outcompete undersea cables. Speed of light in space is ~40% faster than cable and the path is direct, instead of following the outline of land masses."*
>
> ELON MUSK · @ELONMUSK · 2:32PM EST · POST

When a reader multiplied the original numbers out to 100 terabits per second a satellite, Musk replied: "If we shrink laser size, assuming 10X closer proximity, and have small gimbal angle for in-plane links, maybe we can get to 100 lasers/sat."

Orbital compute and the Anthropic deal @elonmusk

Replying to a post that summarized his remarks as a plan to launch StarMind computers into orbit from the third quarter of 2027, Musk wrote: "When we unveil the details of the Starmind computer next year, it will blow people's minds." On Sunday he endorsed a secondhand account, from the Invest Like the Best podcast, that Larry Page puts the share of computing that will end up in orbit at "five nines": "Larry Page has made many accurate predictions and this is one of them." Chamath Palihapitiya argued on the All-In podcast that if data-center supply tightens, SpaceX could end its compute deal with Anthropic on 90 days' notice and bring it back at two to three times the price. Musk replied at 8:58pm EST: "It works both ways, of course. If compute becomes plentiful and cheap, SpaceX revenue will drop. No crying in the casino for either party."

> QUICK HIT
>
> **Qualcomm and Starlink Mobile.** Qualcomm CEO Cristiano Amon called the SpaceX partnership "one of the largest 5G deployments in the United States," with SpaceX building the base stations and Qualcomm supplying the radio silicon. Musk quoted him: "Excited to be working with @Qualcomm on breakthrough technologies for 5G & beyond."

II · FACTORIES AND CHIPS

Cybertruck gives way to Cybercab as Terafab adds a research fab

Musk confirmed that the Cybertruck line at Gigafactory Texas winds down around the middle of next year to make room for Cybercab, with an extra shift to work through orders. He pitched a research chip fab feeding Terafab, and Nvidia's board opened an inquiry into its chief executive.

~8 mo	40 kWh	Nov 18
Cybertruck production Musk said remains before the line shuts down	Cybercab battery pack Musk said coming efficiency gains will allow	Date three House committee chairs asked Jensen Huang to brief them behind closed doors

Elon Musk · Cybertruck @elonmusk

Late Thursday, at 1:30am EST, Musk wrote: "Tesla will be winding down production of the Cybertruck at Gigafactory Texas around the middle of next year to make room for increased Cybercab production," as CNBC reported. On Saturday he added that the order backlog "is currently far more than we can make before dismantling the line for Cybercab expansion around June/July next year. We will add an extra shift to try to get all orders done in time." Early Monday he quoted commentator James Douma ("cars are going to become robots … Get it while you can"):

> *"The Cybertruck order rate far exceeds our production rate. We will try to make as many as possible before shutting down the line in about 8 months, but will have to return deposits for a lot of people. This is not a case of Cybertruck being unprofitable, as has been misreported by legacy media, but rather that the opportunity cost of the complexity of Cybertruck is vastly outweighed by the value of incremental Cybercab production."*
>
> ELON MUSK · @ELONMUSK · 2:20AM EST · POST

A user contrasted a 120 kWh "CyberSUV" with a 48 kWh Cybercab in a robotaxi fleet. Musk replied: "There are efficiency improvements coming that will allow for a 40kWh Cybercab pack. We will make this in liquid silver, so the colors are not too monotonous. The biggest thing most people fail to understand is that there are a very limited number of smart, dedicated people!" Asked why Tesla doesn't simply build more factory space, he said a high-volume factory is extremely hard to build: "Prototypes are trivial by comparison."



Building 1 at Tesla's Gigafactory Texas near Austin, shown on June 17, 2022.
PHOTO: LARRY D. MOORE / WIKIMEDIA COMMONS · CC BY 4.0 (CROPPED)

A research fab for Terafab @elonmusk

A widely shared summary of the Baron interview said Tesla is building a research fab at Giga Texas covering logic and memory fabrication, lithography, advanced packaging and testing, with a goal of cutting a chip development cycle from a year to about a month, while Terafab in Grimes County handles volume production. Musk quoted it at 6:53pm EST: "If you want to iterate fast with radically new chip designs, join the Terafab project!" On a post about wafers spending most of their time waiting between tools, he added: "If the wafer is transported in vacuum, then no pump down is required when it transfers to a machine that requires vacuum. This would greatly improve cycle time and yield."

Nvidia's board opens an inquiry Jensen Huang · @JensenHuang

Nvidia's board said Sunday it is opening an independent investigation into reports that Jensen Huang was the "target of a foreign intelligence effort," the Associated Press reported. In a memo to staff, Huang acknowledged a relationship with Alice Huang, an executive at Megaspeed, a Singapore-based Nvidia customer, and apologized; the board said it "has no reason to believe that either national security or NVIDIA secrets were compromised," CNBC reported. The chairs of the House China, Foreign Affairs and Intelligence committees asked Huang for a closed-door briefing on Nov. 18. Investor Shaun Maguire wrote that Huang "did something incredibly dumb" but is "irreplaceable for America's AI leadership." Musk quoted him at 6:50pm EST: "I agree."

III · MODELS AND AGENTS

Meta brings Muse Wonder's memory to everyone; OpenAI launches GPT-6 Cyber

Meta rolled specialized models trained with Muse Wonder into Muse and Meta AI. OpenAI introduced GPT-6 Cyber, Thinking Machines released Inkling 2 Small, and Musk promised Grok Bot fixes overnight after subscribers hit their limits.

2.4x	67%	31
Reduction in forgetting critical information in Muse, per Meta	Improvement in finishing everyday tasks when unexpected issues arise, per Meta	Inkling 2 Small's score on the Artificial Analysis Intelligence Index, per Thinking Machines

Mark Zuckerberg and Alexandr Wang @finkd · CEO, Meta · @alexandr_wang · Chief AI officer, Meta

Zuckerberg posted at 1:08pm EST:

> *"Upgraded the models for Muse and Meta Al, bringing Active Memory and Active Focus from Muse Wonder to everyone. In Muse, these new specialized models are 2.4x less likely to forget critical info and are 67% better at self correcting and completing everyday tasks which means less repeating yourself and fewer dead ends."*
> MARK ZUCKERBERG · @FINKD · 1:08PM EST · POST

Meta's research blog says the models are trained to tell temporary information from permanent and stale from current; on its Active Memory evaluation, the current Muse model used critical context correctly in 90% of tasks against 76% for the previous one. Wang wrote that "muse and meta ai will run on new specialized models trained using muse wonder." The same day Muse opened Spaces, custom tabs that the assistant builds and keeps updated. Zuckerberg said he made "a BBQ game, an event tracker, and a home control dashboard," and David Singleton built a Thanksgiving planning space to be his "command center on the big day." Wang teased "training everyone's muse to dance before our launch this week."

Greg Brockman · GPT-6 Cyber @gdb · President, OpenAI

OpenAI introduced GPT-6 Cyber on Thursday, saying the model powered the Linux kernel research it had announced the day before and is in use by security teams at companies including Mozilla and Palo Alto Networks. BleepingComputer tallied 81 Firefox flaws alongside the 948 Linux privilege escalations. Brockman wrote at 12:31pm EST: "GPT-6 Cyber: a major advance for security research, already being used to help defend security-critical software such as the Linux kernel, Firefox, and many other projects. We've invested a great deal in

trust, specifically training the model to stay within the scope specified ..."

Mira Murati · Inkling 2 Small @miramurati · CEO, Thinking Machines Lab

Thinking Machines introduced Inkling 2 Small, saying it scores 31 on the Artificial Analysis Intelligence Index and is available for fine-tuning and inference on its Tinker platform, with a full weight release planned. Murati wrote at 6:45pm EST: "Inkling 2 Small is available today on Tinker, with weights coming soon. It's a capable small model. We look forward to seeing what you make with it."

Elon Musk · Grok Bot usage @elonmusk

A subscriber complained that a $200 plan that had lasted seven days "full throttle" the week before now lasted 14 hours under moderate load. Musk replied "I will dig into this today," then posted at 8:17am EST: "Grok Bot fixes for improved token efficiency and smarter behavior are underway." On Sunday night:

> *"Just did a round table review with the Grok @Bot engineering team. Fixes are being implemented overnight and you should notice a dramatic improvement by tomorrow. If not, reply to this post and we will address any remaining issues right away."*
>
> ELON MUSK · @ELONMUSK · 9:25PM EST · POST

Aaron Levie, Greg Brockman and Kevin Weil @levie · CEO, Box · @gdb · @kevinweil

Brockman posted on Sunday at 5:42pm EST: "frontier models, with lots of test time compute, are becoming the engine of scientific discovery." Kevin Weil of OpenAI added: "Yes and—in combination with (AI-driven) physical labs!" Levie quoted Brockman: "We're moving from an era of AI where the output comes roughly at the rate that you can chat back and forth with the AI to work with you, to a period where you can set a problem for a set of agents to go off and do an arbitrarily large amount of work for you. This means productivity gains that are maybe 20-50% to gains that don't really have an upper limit."

IV · THE SCALING ARGUMENT

Sutton says brute force is not enough; Collison asks what winning looks like

A two-sentence post from Richard Sutton drew Chollet and Gu into an argument over what scaling buys. Collison questioned the shape of the upside, Buterin defended slowing data centers, Tim Sweeney and Michael Dell pushed back on the backlash, and Mensch warned French lawmakers about a copyright bill.

7–0	**1 GW**	**Nov 19**
Bloomington, Indiana, council vote to bar superintelligent AI within city limits	Compute capacity Mistral says it is developing, mostly in Europe and largely in France	Date France's National Assembly is due to take up the Darcos bill

Richard Sutton, François Chollet and Albert Gu

@RichardSSutton · @fchollet · @_albertgu

Sutton, the Turing Award-winning author of the 2019 essay "The Bitter Lesson," posted on Saturday at 1:57pm EST: "Brute force computation is not enough for intelligence. Clever new methods are needed. The science of AI is far from finished. IMHO." Chollet quoted him:

> *"Data scaling and compute scaling are *multipliers* for the generalization power of an approach. Trillions of dollars in investment can make very inefficient, low-generalization approaches look adequate, but that doesn't make them the ultimate end-all-be-all of AI. You should work on methods that *can* scale, because if you don't benefit from this multiplier you simply can't compete. But you should keep in mind scaling is only a multiplier: a cleverer approach will scale *better*."*
>
> FRANÇOIS CHOLLET · @FCHOLLET · 5:52PM EST · POST

Gu wrote: "it's hilarious how many have taken the bitter lesson to mean just scaling up transformers. it's about scaling, but innovations in algorithmic primitives can lead to much more efficient scaling." Yann LeCun reposted a 2023 claim that "Altman, Hassabis, and Amodei are the ones doing massive corporate lobbying" with "Reminder from 3 years ago," and defended his joint-embedding predictive architecture (JEPA) research: "There is more time between the transformer paper (2017) and ChatGPT (Late 2022) than between the original JEPA paper (2022) and now."

COMING UP

- **Nov 9** Jim Fan's keynote at the Conference on Robot Learning in Austin, per his post
- **Nov 13** Starship Flight 15, the first attempt to catch both stages, per SpaceX
- **Nov 18** Closed-door briefing House chairs requested from Jensen Huang, per the committee
- **Nov 19** France's National Assembly takes up the Darcos bill, per the Senate dossier

Patrick Collison · two questions @patrickc · CEO, Stripe

Collison posted a long note on Sunday. On the upside case:

> *"People often highlight the risks involved in AI, which are surely very real. I'm struck, however, by how hard it is to understand the contours of the positive scenario. … I'm curious what the larger equilibrium is supposed to look like. Do people work? If so, why? What is their marginal contribution? If not, how are resources allocated? What is scarce? What is the ratio between humans and AIs? Are humans in control, whatever that means? … Given how rapid the rate of progress in AI may be, it feels strange for the upside case to be so opaque."*
>
> PATRICK COLLISON · @PATRICKC · 1:22PM EST · POST

His second point: granting AIs moral worth "would presumably mean that the aggregate worth of AIs would soon outweigh that of humans in any utilitarian calculus."

Vitalik Buterin · the case for slowing down

@VitalikButerin · Co-founder, Ethereum

Buterin wrote at 5:34am EST that he disagreed with "the strong emotional reaction against the anti-data-center movements": "algorithmic progress means that whatever you would be able to do with more and bigger data centers in 2 years, you will be able to do with the data centers we have today in 4 years - probably sooner." He argued that more time helps "harden the world against both already-visible cybersecurity issues and any bio and other risks soon to come," and asked: "Given this, a slowdown in compute hyper-scaling seems quite net good?"

Tim Sweeney and Michael Dell · the backlash

@TimSweeneyEpic · CEO, Epic Games · @MichaelDell · CEO, Dell Technologies

Told that AI "kills nature," Sweeney replied that data centers "are a negligible fraction of developed real estate, use a minimal amount of water (all modern construction uses recirculatory cooling), and are a very useful tool for nature conservation analysis." On Thursday he had written that good-faith critics "are vastly outnumbered by these astroturfed anti-AI influence campaigns." Bloomington's council voted unanimously to bar deploying superintelligent AI in the city, including through browsers and apps. Dell asked: "Curious if they also banned linear algebra, calculus and statistics?" Musk endorsed a clip of his own argument that opposition to data centers is temporary because "Memphis is booming."

Arthur Mensch · France's Darcos bill @arthurmensch · CEO, Mistral AI

In a French-language X article on Sunday, Mensch said the Darcos bill would create a presumption that an AI model used a work whenever there is "a simple clue" (translated), reversing the burden of proof: showing that a work was not used, in a corpus of billions of words, would mean exposing trade secrets. Mistral trained Mistral Large 4 at Bruyères-le-Châtel and is developing 1 gigawatt of capacity, mostly in Europe and largely in France, and he warned that the legal risk "will make the difference for us and our clients" (translated) when choosing data-center sites. He proposed instead a levy on the revenue of every AI provider operating in Europe to fund European creators. The Senate passed the bill unanimously on April 8, and the National Assembly takes it up on Nov. 19. Senator Laure Darcos answered with an open letter the same day, Frandroid reported.

V · BUILDERS' BENCH

Omarchy stays free, Uncle Bob stops reading the code, and Karpathy one-shots a video site

David Heinemeier Hansson said Omarchy will never be monetized and previewed Secure Boot and local AI for 4.6. Uncle Bob Martin described building through agents, and Andrej Karpathy and Clément Delangue showed what weekend and open-source building looks like now.

$20M+	47K	100K
Pledges Heinemeier Hansson said the nonprofit Omacom Foundation has raised	Videos in Karpathy's TristanTube, built with Claude Code in three minutes	Public robotics datasets on Hugging Face, up 48%, per Delangue

David Heinemeier Hansson @dhh · Creator, Ruby on Rails and Omarchy

Heinemeier Hansson posted on Sunday at 9:25am EST:

> *"Omarchy is forever free. There's no monetization scheme coming. I'm already rich and the Omacom Foundation is a well-founded nonprofit. The only ulterior motive here is to make Linux win on the desktop and create the perfect, malleable computer for all to enjoy!"*
>
> DAVID HEINEMEIER HANSSON · @DHH · 9:25AM EST · POST

He said the foundation has "raised over twenty million in pledges. That's enough." Early Monday he announced "Secure boot is coming to Omarchy," so people can install it "without having to fiddle with the bios. (And so OEMs can preinstall it from the factory!)," and said Omarchy 4.6 will recommend "the very best model that's available for your machine" for local AI. He also said he now runs Omarchy on an ASUS Zenbook 14 and 16 as Snapdragon support grows.

Uncle Bob Martin @unclebobmartin · Author, Clean Code

Answering a "BREAKING: Uncle Bob Martin has given up" jab, Martin wrote: "Disengaging from syntax is not disengaging from engineering." On a pre-production project, he said, he makes passes for "Architecture, coverage, race & concurrency, security, API, etc.," each requiring "significant interrogation and review of verification procedures; but not review of the code," with his /ub-mode rules keeping the agents disciplined. He added: "Ours will be the craft of constraining the agents to produce good clean code," and "If six decades in this industry has taught me anything it is that everything changes every year. Keep up or die."

Andrej Karpathy · TristanTube @karpathy

After Tristan, a friend, published his top YouTube channels, Karpathy turned the list into TristanTube: "47K super nerdy videos +TV mode for laid back watching. (One-shotted with Claude Code in 3 minutes)." It follows HackerTube, his earlier hack that filtered 43 million Hacker News comments down to 6,454 recommended videos.

Clément Delangue · open robotics @ClementDelangue · CEO, Hugging Face

Delangue wrote that "open-source is critical for robotics AI" and listed almost 100,000 public robotics datasets (up 48%), over 30,000 robotics models (up 41%), and more than one million downloads of LeRobot, whose version 0.7.0 adds world action models and remote GPU inference.

QUICK HIT

Nvidia's Jim Fan gives a keynote at the Conference on Robot Learning (CoRL) in Austin this afternoon, promising "Real robot. Real time. Uncut shot. Autonomous." Fei-Fei Li replied "Looking forward!" to news that the conference is sold out.

ALSO ON THE WIRE

Paul Graham said a startup he advised to sell to its Y Combinator batchmates "just crossed 10% of the batch using it." Alexandr Wang recalled a 2017 lunch about whether the International Mathematical Olympiad crowd "could be the next paypal mafia." Buterin argued that Ether should be treated as a store of value rather than a fee-burning token. Tobi Lütke pointed to Alexis Ohanian's agent buying more from Shopify stores.

Takeaways

1. Flight 15 is set up as Starship's biggest test of reuse yet: both stages are to come home to Starbase's tower, no earlier than Friday.
2. Tesla is ending a truck line Musk says is not unprofitable, because the factory space and people are worth more as Cybercab capacity.
3. Musk's orbital pitch now runs from 10-terabit laser links to orbital computers, and he concedes the economics swing both ways for compute customers.
4. Memory and persistence are the new consumer-assistant selling points: Meta shipped Muse Wonder's to everyone, while Grok Bot's usage limits drew fixes from Musk.
5. The scaling consensus is fraying at the edges: Sutton, Chollet and Gu want smarter methods, and Buterin welcomes a slower compute buildout.
6. Rules are arriving locally and nationally, from Bloomington's superintelligence ban to France's Darcos bill, and builders are pushing back.

On the beat

POSTED · 34

Elon Musk, Uncle Bob Martin, Andrej Karpathy, John Carmack, Patrick Collison, Yann LeCun, Palmer Luckey, Vitalik Buterin, Tobi Lütke, David Heinemeier Hansson, Jeff Dean, François Chollet, Fei-Fei Li, Greg Brockman, Mustafa Suleyman, Paul Graham, Clément Delangue, Jim Fan, Alexandr Wang, Mira Murati, Mark Zuckerberg, Noam Brown, David Singleton, Joe Gebbia, Guillermo Rauch, Arthur Mensch, Mark Chen, Michael Dell, Aaron Levie, Tim Sweeney, Kevin Weil, Aidan Gomez, Suhail Doshi, Albert Gu.

SILENT · 16

Marc Andreessen, Jensen Huang, Ashok Elluswamy, Sam Altman, Demis Hassabis, Dario Amodei, Satya Nadella, Andrew Ng, Michael Truell, Tom Brown, Ilya Sutskever, Lisa Su, Sundar Pichai, Kanjun Qiu, Nat Friedman, Guillaume Lample.

Added Richard Sutton (@RichardSSutton), Turing Award-winning reinforcement-learning researcher at the University of Alberta. The list now stands at 51.

The SI Sentinel is compiled from public posts on X and published reporting. Quotes are reproduced verbatim unless marked as translated or paraphrased; emoji in quoted posts are omitted.